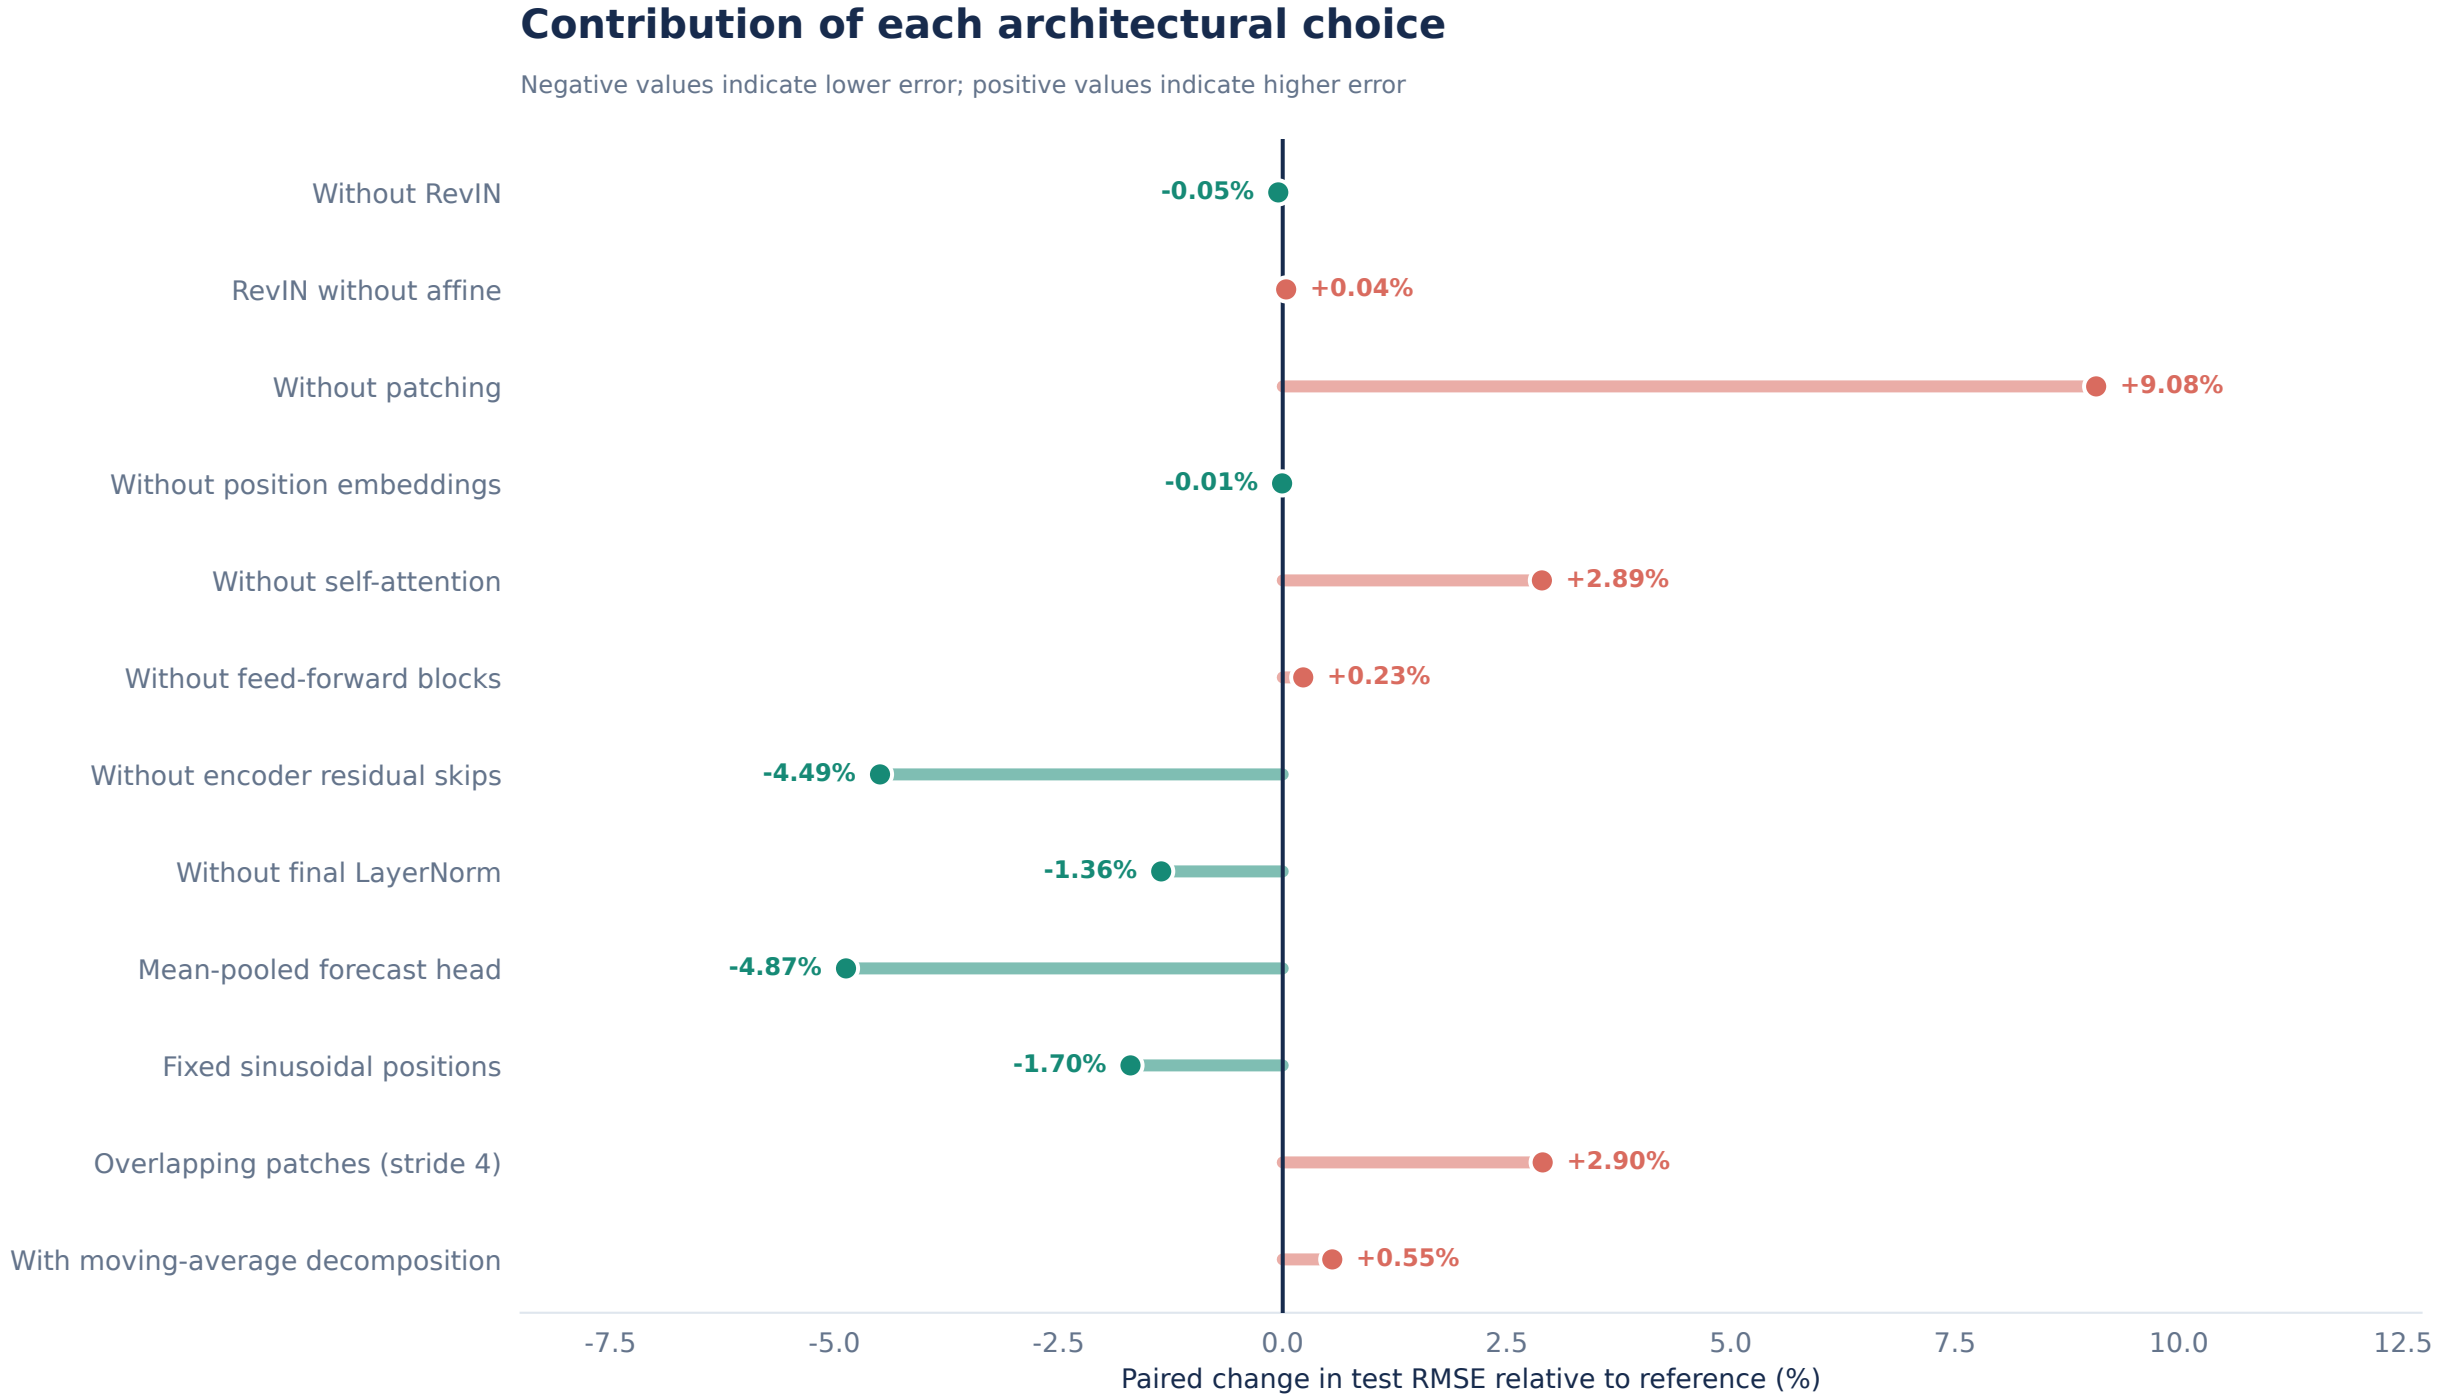
Contribution of each architectural choice
Negative values indicate lower error; positive values indicate higher error
Without RevIN
-0.05%
RevIN without affine
+0.04%
Without patching
+9.08%
Without position embeddings
-0.01%
Without self-attention
+2.89%
Without feed-forward blocks
+0.23%
Without encoder residual skips
-4.49%
Without final LayerNorm
-1.36%
Mean-pooled forecast head
-4.87%
Fixed sinusoidal positions
-1.70%
Overlapping patches (stride 4)
+2.90%
With moving-average decomposition
+0.55%
-7.5
-5.0
-2.5
0.0
2.5
5.0
7.5
10.0
12.5
Paired change in test RMSE relative to reference (%)